

Elements Of Data Science - F2020

Week 3: Pandas, Data Exploration and Visualization

9/28/2020

TODOs

- Read Practical Statistics for Data Scientists, Chapter 3 [EBSCO](#)
- (Optional) Data Science From Scratch, Chapter 5,6,7 [EBSCO](#)
- (Optional) Seaborn Tutorial <https://seaborn.pydata.org/tutorial.html>
- Complete Week 3 Quiz
- HW1 next week, after Hypothesis Testing

Questions?

TODAY

- Pandas
- Data Exploration
- Visualization in Python

Intro to Pandas



Pandas is an open source, BSD-licensed library providing:

- high-performance, easy-to-use data structures and
- data analysis tools

```
In [2]: # usually imported as pd  
import pandas as pd
```

- **Series:** 1D array with a flexible index
- **Dataframe:** 2D matrix with flexible index and column names

Pandas Series

- 1D array of data (any numpy datatype) plus an associated **index** array

```
In [3]: s = pd.Series(np.random.rand(4))  
s
```

```
Out[3]: 0    0.045003  
        1    0.500485  
        2    0.502491  
        3    0.577879  
        dtype: float64
```

```
In [4]: # return the values of the series  
s.values
```

```
Out[4]: array([0.04500301, 0.50048463, 0.50249066, 0.57787909])
```

```
In [5]: # return the index of the series  
s.index
```

```
Out[5]: RangeIndex(start=0, stop=4, step=1)
```

Pandas Series Cont.

- index is flexible, can be anything hashable (integers, strings, ...)

```
In [6]: # create Series from array and set index
s = pd.Series(np.random.rand(3), index=['a', 'b', 'c'])
s
```

```
Out[6]: a    0.783338
        b    0.634322
        c    0.739958
        dtype: float64
```

```
In [7]: s['a']
```

```
Out[7]: 0.7833376791406063
```

```
In [8]: s[['b', 'c']]
```

```
Out[8]: b    0.634322
        c    0.739958
        dtype: float64
```

Pandas Series Cont.

```
In [9]: # Can create series with index from a dictionary
s = pd.Series({'a':1, 'b':2, 'c':3, 'd':4})
s
```

```
Out[9]: a    1
        b    2
        c    3
        d    4
        dtype: int64
```

```
In [10]: s[s.index[-2:]]
```

```
Out[10]: c    3
         d    4
         dtype: int64
```


Pandas DataFrame

- tabular datastructure
- each column a single datatype
- contains both row and column indices
- single column == Series

Pandas DataFrame Cont.

```
In [11]: df = pd.DataFrame({'Year':[2017,2018,2018,2019],  
                             'Class_Name':['A','A','B','A'],  
                             'Measure1':[2.1,3.0,2.4,1.9]  
                             })
```

```
In [12]: df
```

Out[12]:

	Year	Class_Name	Measure1
0	2017	A	2.1
1	2018	A	3.0
2	2018	B	2.4
3	2019	A	1.9

```
In [13]: print(df)
```

```
   Year Class_Name Measure1  
0  2017          A        2.1  
1  2018          A        3.0  
2  2018          B        2.4  
3  2019          A        1.9
```

```
In [14]: display(df)
```

	Year	Class_Name	Measure1
0	2017	A	2.1
1	2018	A	3.0
2	2018	B	2.4
3	2019	A	1.9

Pandas DataFrame Cont.

```
In [15]: data = [[2017, 'A', 2.1],  
                 [2018, 'A', 3.0],  
                 [2018, 'B', 2.4],  
                 [2019, 'A', 1.9]]
```

```
In [16]: df = pd.DataFrame(data,  
                           columns=['Year', 'Class_Name', 'Measure1'],  
                           index=['001', '002', '003', '004'])  
  
df.shape
```

Out[16]: (4, 3)

```
In [17]: df
```

Out[17]:

	Year	Class_Name	Measure1
001	2017	A	2.1
002	2018	A	3.0
003	2018	B	2.4
004	2019	A	1.9

Pandas Attributes

- Get shape of DataFrame : `shape`

```
In [18]: df.shape # rows, columns
```

```
Out[18]: (4, 3)
```

- Get index values : `index`

```
In [19]: df.index
```

```
Out[19]: Index(['001', '002', '003', '004'], dtype='object')
```

- Get column values : `columns`

```
In [20]: df.columns
```

```
Out[20]: Index(['Year', 'Class_Name', 'Measure1'], dtype='object')
```

Pandas Indexing/Selection

Select by label:

- `.loc[]`

```
In [21]: df.loc['001']
```

```
Out[21]: Year          2017  
         Class_Name      A  
         Measure1      2.1  
         Name: 001, dtype: object
```

```
In [22]: df.loc['001', 'Measure1']
```

```
Out[22]: 2.1
```

Pandas Indexing/Selection Cont.

Select by position:

- `.iloc[]`

```
In [23]: df.iloc[0]
```

```
Out[23]: Year      2017  
         Class_Name  A  
         Measure1   2.1  
         Name: 001, dtype: object
```

```
In [24]: df.iloc[0,2]
```

```
Out[24]: 2.1
```

Pandas Indexing/Selection Cont.

Selecting multiple rows/columns: use list (fancy indexing)

```
In [25]: df.loc[['002', '004']]
```

Out[25]:

	Year	Class_Name	Measure1
002	2018	A	3.0
004	2019	A	1.9

```
In [26]: df.loc[['002', '004'], ['Year', 'Measure1']]
```

Out[26]:

	Year	Measure1
002	2018	3.0
004	2019	1.9

Pandas Slicing

```
In [27]: # Get last two rows  
df.iloc[-2:]
```

Out[27]:

	Year	Class_Name	Measure1
003	2018	B	2.4
004	2019	A	1.9

```
In [28]: # Get first two rows and first two columns  
df.iloc[:2,:2]
```

Out[28]:

	Year	Class_Name
001	2017	A
002	2018	A

NOTE: `.iloc` is exclusive (start:end+1)

Pandas Slicing Cont.

Can also slice using labels:

```
In [29]: df.loc['002':'004']
```

Out[29]:

	Year	Class_Name	Measure1
002	2018	A	3.0
003	2018	B	2.4
004	2019	A	1.9

```
In [30]: df.loc['002':'004', : 'Class_Name']
```

Out[30]:

	Year	Class_Name
002	2018	A
003	2018	B
004	2019	A

NOTE: `.loc` is inclusive

Pandas Slicing Cont.

How to indicate all rows or all columns? :

```
In [31]: df.loc[:, 'Measure1']
```

```
Out[31]: 001    2.1  
         002    3.0  
         003    2.4  
         004    1.9  
         Name: Measure1, dtype: float64
```

```
In [32]: df.iloc[2:,:] 
```

```
Out[32]:
```

	Year	Class_Name	Measure1
003	2018	B	2.4
004	2019	A	1.9

Pandas Indexing Cont.

Shortcut for indexing:

```
In [33]: df['Class_Name']
```

```
Out[33]: 001    A
         002    A
         003    B
         004    A
         Name: Class_Name, dtype: object
```

```
In [34]: # can use dot notation if there is no space in label
         df.Class_Name
```

```
Out[34]: 001    A
         002    A
         003    B
         004    A
         Name: Class_Name, dtype: object
```

Panda Selection Chaining

Get 'Year' and 'Measure1' for first 3 rows:

```
In [35]: df.iloc[:3].loc[:, ['Year', 'Measure1']]
```

Out[35]:

	Year	Measure1
001	2017	2.1
002	2018	3.0
003	2018	2.4

For records '001' and '003' get last two columns

```
In [36]: df.loc[['001', '003']].iloc[:, -2:]
```

Out[36]:

	Class_Name	Measure1
001	A	2.1
003	B	2.4

Panda Selection Chaining Cont.

For record '001' get last two columns?:

```
In [37]: %xmode minimal
df.loc['001'].iloc[:, -2:] # row with label '001', then all rows, last two columns?
```

Exception reporting mode: Minimal

IndexingError: Too many indexers

```
In [38]: df.loc['001']
```

```
Out[38]: Year          2017
Class_Name          A
Measure1           2.1
Name: 001, dtype: object
```

```
In [39]: df.loc['001'].iloc[-2:] # row with label '001', last two elements of Series
```

```
Out[39]: Class_Name          A
Measure1           2.1
Name: 001, dtype: object
```

Pandas **head** and **tail**

Get a quick view of the first or last rows in a DataFrame

```
In [40]: df.head() # first 5 rows by default
```

Out[40]:

	Year	Class_Name	Measure1
001	2017	A	2.1
002	2018	A	3.0
003	2018	B	2.4
004	2019	A	1.9

```
In [41]: df.tail(2) # only print 2 rows
```

Out[41]:

	Year	Class_Name	Measure1
003	2018	B	2.4
004	2019	A	1.9

Pandas Boolean Mask

```
In [42]: # Which rows have Class_Name of 'A'?  
df.Class_Name == 'A'
```

```
Out[42]: 001      True  
         002      True  
         003     False  
         004      True  
         Name: Class_Name, dtype: bool
```

```
In [43]: # Get all data for rows with with Class_Name 'A'  
df.loc[df.Class_Name == 'A']
```

```
Out[43]:
```

	Year	Class_Name	Measure1
001	2017	A	2.1
002	2018	A	3.0
004	2019	A	1.9

```
In [44]: # Get Measure1 for all records for Class_Name 'A'  
df.loc[df.Class_Name == 'A', 'Measure1']
```

```
Out[44]: 001      2.1  
         002      3.0  
         004      1.9  
         Name: Measure1, dtype: float64
```

Pandas Boolean Mask Cont.

Get all records for class 'A' before 2019

```
In [45]: df.loc[(df.Class_Name == 'A') & (df.Year < 2019)]
```

Out[45]:

	Year	Class_Name	Measure1
001	2017	A	2.1
002	2018	A	3.0

Get all records in a set of years:

```
In [46]: df.loc[df.Year.isin([2017, 2019])]
```

Out[46]:

	Year	Class_Name	Measure1
001	2017	A	2.1
004	2019	A	1.9

Pandas Selection Review

- `.loc[]`
- `.iloc[]`
- Fancy Indexing
- Slicing
- Chaining
- `head` and `tail`
- Boolean Mask

Pandas Sorting

```
In [47]: df.sort_values(by=['Measure1']).head(3)
```

Out[47]:

	Year	Class_Name	Measure1
004	2019	A	1.9
001	2017	A	2.1
003	2018	B	2.4

```
In [48]: df.sort_values(by=['Measure1'], ascending=False).head(3)
```

Out[48]:

	Year	Class_Name	Measure1
002	2018	A	3.0
003	2018	B	2.4
001	2017	A	2.1

```
In [49]: df.sort_values(by=['Year', 'Measure1']).head(3)
```

Out[49]:

	Year	Class_Name	Measure1
001	2017	A	2.1
003	2018	B	2.4
002	2018	A	3.0

Questions?

Data Exploration and Visualization

For a new set of data, would like to know:

- amount of data (rows, columns)
- range (min, max)
- counts of discrete values
- central tendencies (mean, median)
- dispersion or spread (variance, IQR)
- skew
- covariance and correlation ...

Yellowcab Dataset

- Records of Yellowcab Taxi trips from January 2017
- more info: <https://www1.nyc.gov/site/tlc/about/tlc-trip-record-data.page>

Loading Datasets from CSV (Comma Separated Values)

- columns separated by delimiter, eg. comma, tab (\t), pipe (|)
- one row per record, observation
- often, strings quoted
- often, first row contains column headings
- often, comment rows starting with #

```
In [50]: !head ../data/yellowcab_demo.csv
```

```
pickup_datetime,dropoff_datetime,trip_distance,fare_amount,tip_amount,payment_type
2017-01-05 14:49:04,2017-01-05 14:53:53,0.89,5.5,1.26,Credit card
2017-01-15 01:07:22,2017-01-15 01:26:47,2.7,14.0,0.0,Cash
2017-01-29 09:55:00,2017-01-29 10:04:43,1.41,8.0,0.0,Cash
2017-01-10 05:40:12,2017-01-10 05:42:22,0.4,4.0,0.0,Cash
2017-01-06 17:02:48,2017-01-06 17:16:10,2.3,11.0,0.0,Cash
2017-01-14 19:03:14,2017-01-14 19:08:41,0.8,5.5,,Credit card
2017-01-06 18:51:52,2017-01-06 18:55:45,0.2,4.5,0.0,Cash
2017-01-04 20:47:30,2017-01-04 21:01:24,2.68,11.5,,Credit card
2017-01-21 09:44:28,2017-01-21 09:48:13,0.6,4.5,0.0,Cash
```

Loading Datasets with Pandas

```
In [51]: import pandas as pd
df = pd.read_csv('../data/yellowcab_demo.csv',
                 sep=',',
                 header=0,
                 parse_dates=['pickup_datetime', 'dropoff_datetime'])
```

```
In [52]: # display first 5 rows
df.head(5)
```

Out[52]:

	pickup_datetime	dropoff_datetime	trip_distance	fare_amount	tip_amount	payment_type
0	2017-01-05 14:49:04	2017-01-05 14:53:53	0.89	5.5	1.26	Credit card
1	2017-01-15 01:07:22	2017-01-15 01:26:47	2.70	14.0	0.00	Cash
2	2017-01-29 09:55:00	2017-01-29 10:04:43	1.41	8.0	0.00	Cash
3	2017-01-10 05:40:12	2017-01-10 05:42:22	0.40	4.0	0.00	Cash
4	2017-01-06 17:02:48	2017-01-06 17:16:10	2.30	11.0	0.00	Cash

Get Size of Dataset

```
In [53]: df.shape
```

```
Out[53]: (1000, 6)
```

```
In [54]: # rows  
f'{df.shape[0]} rows'  
'1000 rows'
```

```
Out[54]: '1000 rows'
```

```
In [55]: f'{df.shape[1]} columns'  
'6 columns'
```

```
Out[55]: '6 columns'
```

```
In [56]: 'number of rows: {}, number of columns: {}'.format(*df.shape)
```

```
Out[56]: 'number of rows: 1000, number of columns: 6'
```


What are the column names?

```
In [57]: df.columns
```

```
Out[57]: Index(['pickup_datetime', 'dropoff_datetime', 'trip_distance', 'fare_amount',  
              'tip_amount', 'payment_type'],  
              dtype='object')
```

```
In [58]: df.columns.values
```

```
Out[58]: array(['pickup_datetime', 'dropoff_datetime', 'trip_distance',  
              'fare_amount', 'tip_amount', 'payment_type'], dtype=object)
```

```
In [59]: df.columns.tolist()
```

```
Out[59]: ['pickup_datetime',  
          'dropoff_datetime',  
          'trip_distance',  
          'fare_amount',  
          'tip_amount',  
          'payment_type']
```

What are the column datatypes?

```
In [60]: df.dtypes
```

```
Out[60]: pickup_datetime    datetime64[ns]  
dropoff_datetime          datetime64[ns]  
trip_distance              float64  
fare_amount                float64  
tip_amount                 float64  
payment_type               object  
dtype: object
```

```
In [61]: type(df.dtypes)
```

```
Out[61]: pandas.core.series.Series
```

Get Summary Info for DataFrame

```
In [62]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   pickup_datetime       1000 non-null   datetime64[ns]
1   dropoff_datetime      1000 non-null   datetime64[ns]
2   trip_distance         1000 non-null   float64
3   fare_amount           1000 non-null   float64
4   tip_amount            910 non-null    float64
5   payment_type          1000 non-null   object  
dtypes: datetime64[ns](2), float64(3), object(1)
memory usage: 47.0+ KB
```

- number of rows
- number of columns
- column names, number of filled values, datatypes
- number of each datatype seen
- size of dataset in memory

Variable (Observation) Types

- **Numeric** (eg. weight, temperature)
 - usually has a zero value
 - describes magnitude
- **Categorical** (eg. class, variety)
 - usually a finite set
 - no order
- **Ordinal** (eg. Likert scale, education level, etc.)
 - usually a finite set
 - has order
 - usually missing zero
 - difference between levels may not be the same

Numeric: Data Ranges

```
In [63]: df.trip_distance.min()
```

```
Out[63]: 0.0
```

```
In [64]: df.trip_distance.max()
```

```
Out[64]: 32.77
```

```
In [65]: df.min(numeric_only=True)
```

```
Out[65]: trip_distance    0.0  
fare_amount           2.5  
tip_amount            0.0  
dtype: float64
```

```
In [66]: df.max(numeric_only=True)
```

```
Out[66]: trip_distance    32.77  
fare_amount           88.00  
tip_amount            22.70  
dtype: float64
```

Numeric: Central Tendency with Mean

- Sample Mean

$$\bar{x} = \frac{1}{n} \sum x_i$$

```
In [67]: df.fare_amount.mean()
```

```
Out[67]: 12.4426
```

```
In [68]: print(f'{df.fare_amount.mean():0.2f}')  
  
df.fare_amount.mean()=12.44
```

- Mean is sensitive to *outliers*
- **Outlier:** a data point that differs significantly from other observations
 - data error
 - effect of heavy tailed distribution?

Numeric: Central Tendency with Median

- Median
 - Divides sorted dataset into two equal sizes
 - 50% of the data is less than or equal to the median

```
In [69]: df.fare_amount.median()
```

```
Out[69]: 9.0
```

- Median is *robust* to outliers
- **Robust:** Not affected by outliers

Numeric: Quantiles/Percentiles

- **Quantile**:: cut point for splitting distribution
- **Percentile**: $x\%$ of data is less than or equal to the x th percentile

```
In [70]: df.fare_amount.quantile(.95) # 95% of the data is less than or equal to x?
```

```
Out[70]: 33.5
```

```
In [71]: df.fare_amount.quantile([.05,.95]) # 90% of the data is between 4 and 33.5
```

```
Out[71]: 0.05    4.0  
        0.95   33.5  
        Name: fare_amount, dtype: float64
```

```
In [72]: df.fare_amount.quantile([0,.25,.5,.75,1]) # Quartiles: 25% of data is between each pair
```

```
Out[72]: 0.00    2.5  
        0.25    6.5  
        0.50    9.0  
        0.75   14.0  
        1.00   88.0  
        Name: fare_amount, dtype: float64
```


Numeric: Spread with Variance

- Sample Variance

$$s^2 = \frac{\sum (x - \bar{x})^2}{n-1}$$

```
In [73]: df.fare_amount.var()
```

```
Out[73]: 116.80859383383383
```

but this is in dollars²!

Numeric: Spread with Standard Deviation

- Sample Standard Deviation

$$s = \sqrt{\frac{\sum (x - \bar{x})^2}{n-1}}$$

```
In [74]: df.fare_amount.std()
```

```
Out[74]: 10.807802451647321
```

- Back in original scale of dollars
- Sensitive to outliers

Numeric: Exploring Spread with IQR

- Quartiles
 - ~25% of data is \leq first quartile, 25th percentile
 - ~50% of data is \leq second quartile, 50th percentile (Median)
 - ~75% of data is \leq third quartile, 75th percentile
- Can find quartiles with: pandas quantile or numpy percentile
- Interquartile Range (IQR)
 - (third quartile - first quartile) or (75th percentile - 25th percentile)

```
In [75]: df.fare_amount.quantile(.75) - df.fare_amount.quantile(.25)
```

```
Out[75]: 7.5
```

- IQR is robust to outliers

Numeric Summary Stats with `.describe`

In [76]: `df.describe()`

Out[76]:

	trip_distance	fare_amount	tip_amount
count	1000.000000	1000.000000	910.000000
mean	2.880010	12.442600	1.766275
std	3.678534	10.807802	2.315507
min	0.000000	2.500000	0.000000
25%	0.950000	6.500000	0.000000
50%	1.565000	9.000000	1.350000
75%	3.100000	14.000000	2.460000
max	32.770000	88.000000	22.700000

Numeric: Exploring Distribution with Skew

- Skewness
 - measures asymmetry of distribution around mean
 - indicates tail to left (neg) or right (pos)
 - skew will lead to difference between median and mean

```
In [77]: df.fare_amount.skew()
```

```
Out[77]: 2.882730031010152
```

Easier to understand with a plot...


Matplotlib.pyplot

```
In [78]: import matplotlib.pyplot as plt
```


```
%matplotlib inline
```

```
In [79]: plt.scatter(df.trip_distance,df.fare_amount)
```

```
Out[79]: <matplotlib.collections.PathCollection at 0x7f1df5b0bf40>
```



```
In [80]: plt.scatter(df.trip_distance,df.fare_amount);
```



Matplotlib Axes

Matplotlib Axes

```
In [81]: fig = plt.figure(figsize=(6,4))


ax = fig.gca()

ax.scatter(x=df.trip_distance,
          y=df.fare_amount,
          marker='x',
          color='red'
          )

ax.set_xlabel('trip_distance')
ax.set_ylabel('fare_amount')

ax.set_xlim([-10,50])
ax.set_ylim([-10,100])

ax.set_title('trip_distance vs fare_amount');
```



Matplotlib: Subplots, Figure and Axis

Matplotlib: Subplots, Figure and Axis

```
In [82]: fig, ax = plt.subplots(1, 2, figsize=(12, 4))


ax[0].scatter(df.trip_distance, df.fare_amount, marker='x', color='blue')
ax[1].scatter(df.trip_distance, df.tip_amount, color='red');

ax[0].set_xlabel('trip_distance')
ax[1].set_xlabel('trip_distance')

ax[0].set_ylabel('fare_amount'), ax[1].set_ylabel('tip_amount')

ax[0].set_title('trip_distance vs fare_amount')
ax[1].set_title('trip_distance vs tip_amount')


fig.suptitle('Yellowcab Taxi Features');
```



Plotting via Pandas

Plotting via Pandas


```
In [83]: ax = df.plot.scatter(x='trip_distance',y='fare_amount');  
ax.set_title('trip_distance vs fare_amount');
```



Univariate Distribution: Histogram


Univariate Distribution: Histogram

```
In [84]: df.fare_amount.plot.hist();
```




Univariate Distribution: Histogram

```
In [84]: df.fare_amount.plot.hist();
```



```
In [85]: ax = df.fare_amount.plot.hist(bins=100)  
ax.set_xlabel('fare_amount');
```



Univariate Distribution: Histogram

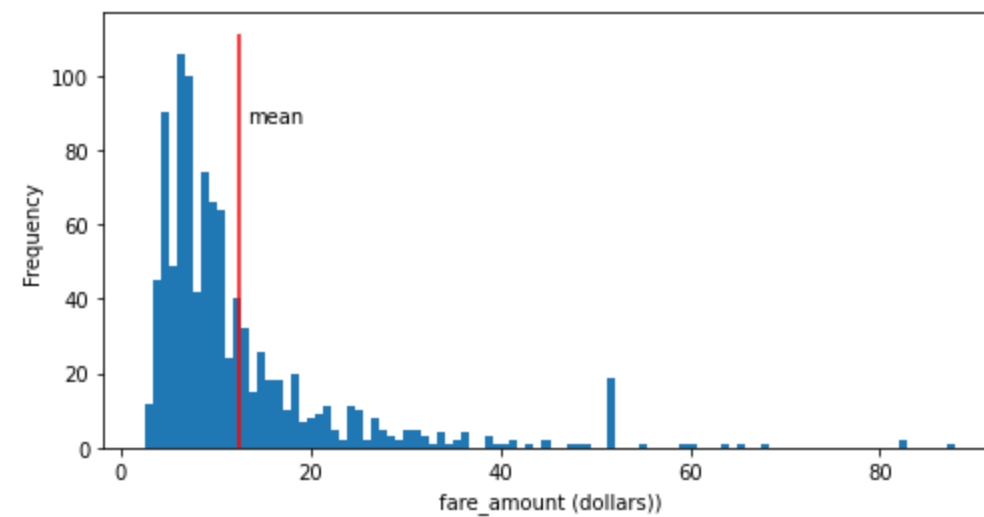
Univariate Distribution: Histogram

```
In [86]: fig,ax = plt.subplots(1,1,figsize=(8,4));

df.fare_amount.plot.hist(bins=100, ax=ax);
ax.set_xlabel('fare_amount (dollars)');

# add a vertical line
ax.vlines(df.fare_amount.mean(), *ax.get_ylim(),color='r');

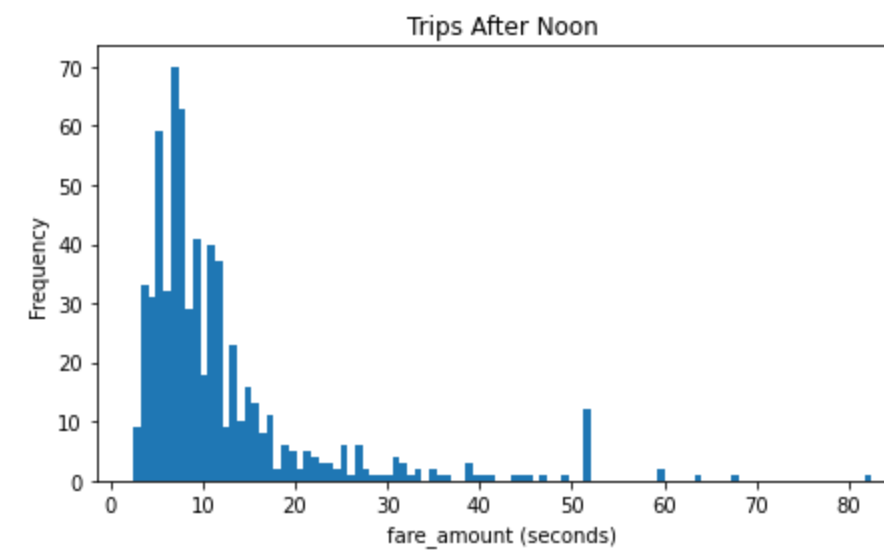
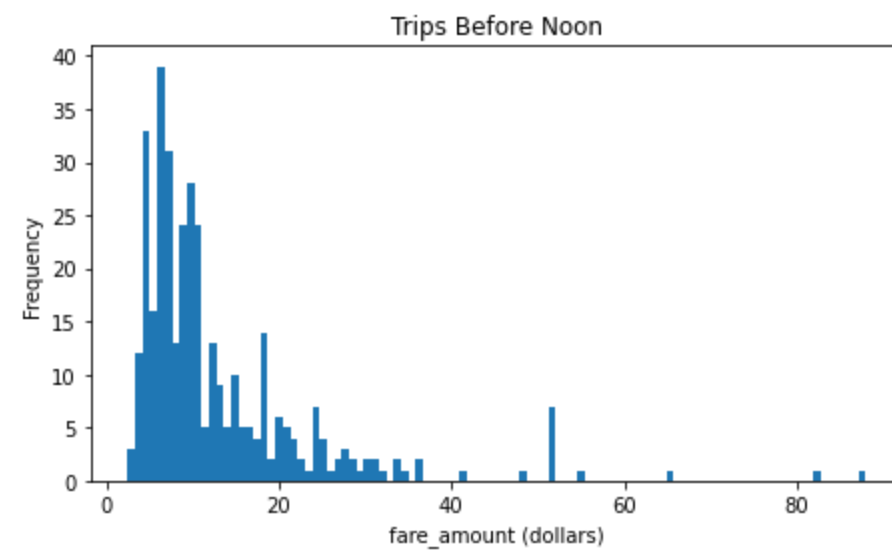
# add some text
ax.text(df.fare_amount.mean()+1,ax.get_ylim()[1]*.75, 'mean');
```



Subplots with Pandas

Subplots with Pandas

```
In [87]: fig,ax = plt.subplots(1,2,figsize=(16,4))
df[df.pickup_datetime.dt.hour < 12].fare_amount.plot.hist(bins=100,ax=ax[0]);
ax[0].set_xlabel('fare_amount (dollars)');
ax[0].set_title('Trips Before Noon');
df[df.pickup_datetime.dt.hour >= 12].fare_amount.plot.hist(bins=100,ax=ax[1]);
ax[1].set_xlabel('fare_amount (seconds)');
ax[1].set_title('Trips After Noon');
```





Sharing Axes

Sharing Axes

```
In [88]: fig, ax = plt.subplots(1, 2, figsize=(16, 4), sharey=True)


df[df.pickup_datetime.dt.hour < 12].fare_amount.plot.hist(bins=100, ax=ax[0]);
ax[0].set_xlabel('fare_amount (dollars)');
ax[0].set_title('Trips Before Noon');
df[df.pickup_datetime.dt.hour >= 12].fare_amount.plot.hist(bins=100, ax=ax[1]);
ax[1].set_xlabel('fare_amount (seconds)');
ax[1].set_title('Trips After Noon');
```



Plotting with Seaborn


Plotting with Seaborn

- Python data visualization library
- Based on matplotlib.
- It provides a high-level interface for drawing attractive and informative statistical graphics.



Plotting with Seaborn

- Python data visualization library
- Based on matplotlib.
- It provides a high-level interface for drawing attractive and informative statistical graphics.




NOTE!!!: Upgrade to 0.11 as of September 2020

```
$ conda install -n eods-f20 seaborn
```


Plotting with Seaborn

- Python data visualization library
- Based on matplotlib.
- It provides a high-level interface for drawing attractive and informative statistical graphics.



NOTE!!!: Upgrade to 0.11 as of September 2020

```
$ conda install -n eods-f20 seaborn
```


```
In [89]: import seaborn as sns  
sns.__version__
```

```
Out[89]: '0.11.0'
```

Univariate Distribution with Seaborn Histplot

Univariate Distribution with Seaborn Histplot


```
In [90]: fig, ax = plt.subplots(1, 1, figsize=(12, 8))  
  
sns.histplot(df.fare_amount, ax=ax);
```



Univariate Distribution: Histogram with KDE and Rugplot


Univariate Distribution: Histogram with KDE and Rugplot

```
In [91]: fig, ax = plt.subplots(1, 1, figsize=(12, 8))  
  
sns.histplot(df.fare_amount, kde=True, ax=ax);  
sns.rugplot(df.fare_amount);
```



Aside: KDE


Aside: KDE



Seaborn Styles


Seaborn Styles

```
In [92]: # for a single plot using a context
with sns.axes_style('whitegrid'):
    sns.histplot(df.fare_amount);
```



Seaborn Styles


```
In [92]: # for a single plot using a context
with sns.axes_style('whitegrid'):
    sns.histplot(df.fare_amount);
```



```
In [93]: # set style globally
sns.set_style('darkgrid')
```


Seaborn Styles

```
In [92]: # for a single plot using a context
with sns.axes_style('whitegrid'):
    sns.histplot(df.fare_amount);
```



```
In [93]: # set style globally
sns.set_style('darkgrid')
```


```
In [94]: sns.histplot(x=df.fare_amount);
```



Univariate Distributions: Boxplot

Univariate Distributions: Boxplot


```
In [95]: fig, ax = plt.subplots(1, 1, figsize=(12, 8))  
sns.boxplot(x=df.fare_amount, ax=ax);
```



Univariate Distributions: Boxplot


Univariate Distributions: Boxplot

```
In [96]: fig, ax = plt.subplots(1, 1, figsize=(6, 4))  
sns.boxplot(x=df.fare_amount, ax=ax);
```



Univariate Distributions: Boxplot

```
In [96]: fig, ax = plt.subplots(1, 1, figsize=(6, 4))  
sns.boxplot(x=df.fare_amount, ax=ax);
```





- first quartile
- second quartile (Median)
- third quartile
- whiskers (usually $1.5 \times \text{IQR}$)
- outliers

Combining Plots with Subplots

Combining Plots with Subplots

```
In [97]: fig, ax = plt.subplots(1, 2, figsize=(16, 6))  
  
sns.histplot(x=df.fare_amount, ax=ax[0]);  
sns.boxplot(x=df.fare_amount, ax=ax[1]);
```




Other Univariate Distribution Visualizations

Other Univariate Distribution Visualizations

```
In [98]: fig, ax = plt.subplots(1, 3, figsize=(16, 4))

sns.stripplot(x='fare_amount', data=df[:200], ax=ax[0])
sns.violinplot(x='fare_amount', data=df, ax=ax[1])
sns.swarmplot(x='fare_amount', data=df[:200], ax=ax[2]);
```



Bivariate: Evaluating Correlation

Bivariate: Evaluating Correlation

- **Correlation:** the degree to which two variables are linearly related

- Pearson Correlation Coefficient: $\rho_{XY} = \frac{cov(X,Y)}{\sigma_X \sigma_Y}$

- Sample Correlation: $r = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{(n-1)s_x s_y}$

- Takes values between:
 - -1 (highly negatively correlated)
 - 0 (not correlated)
 - 1 (highly positively correlated)

Bivariate: Evaluating Correlation

- **Correlation:** the degree to which two variables are linearly related

- Pearson Correlation Coefficient: $\rho_{XY} = \frac{cov(X,Y)}{\sigma_X \sigma_Y}$

- Sample Correlation: $r = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{(n-1)s_x s_y}$

- Takes values between:
 - -1 (highly negatively correlated)
 - 0 (not correlated)
 - 1 (highly positively correlated)

```
In [99]: df.trip_distance.corr(df.fare_amount)
```

Bivariate: Evaluating Correlation

- **Correlation:** the degree to which two variables are linearly related

- Pearson Correlation Coefficient: $\rho_{XY} = \frac{cov(X,Y)}{\sigma_X \sigma_Y}$


- Sample Correlation: $r = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{(n-1)s_x s_y}$

- Takes values between:
 - -1 (highly negatively correlated)
 - 0 (not correlated)
 - 1 (highly positively correlated)

```
In [99]: df.trip_distance.corr(df.fare_amount)
```


Pearson Correlation


Pearson Correlation



Bivariate: Scatterplot


Bivariate: Scatterplot

```
In [101]: sns.scatterplot(x='trip_distance',y='fare_amount',data=df);
```




Bivariate: Scatterplot

```
In [101]: sns.scatterplot(x='trip_distance',y='fare_amount',data=df);
```




```
In [102]: sns.scatterplot(x='trip_distance',y='fare_amount',data=df,alpha=0.2);
```



Bivariate: Add Regression Line

Bivariate: Add Regression Line


```
In [103]: fig, ax = plt.subplots(1, 1, figsize=(12, 8))  
  
sns.regplot(x='trip_distance', y='fare_amount', data=df, ax=ax);
```



Bivariate: Joint Plot

Bivariate: Joint Plot

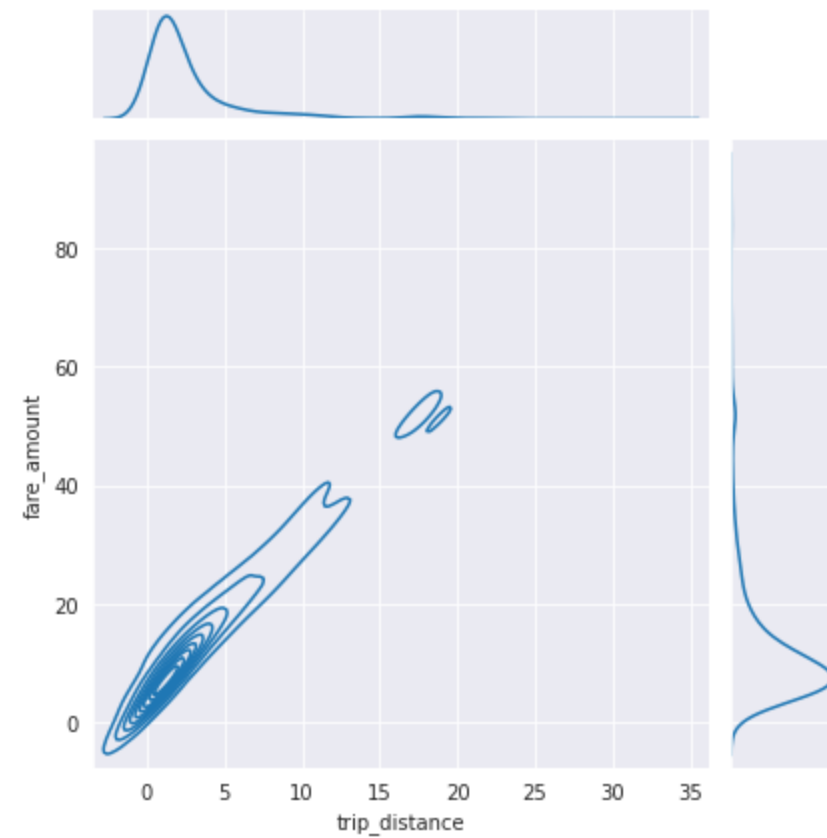
```
In [104]: sns.jointplot(x='trip_distance',y='fare_amount',data=df,alpha=0.5);
```



Bivariate: Joint Plot with KDE

Bivariate: Joint Plot with KDE


```
In [105]: sns.jointplot(x='trip_distance', y='fare_amount',  
                        data=df,  
                        kind='kde');
```



Comparing Multiple Variables with `pairplot`

Comparing Multiple Variables with `pairplot`

```
In [106]: sns.pairplot(df[['trip_distance', 'fare_amount', 'tip_amount']]);
```



Categorical Variables: Frequency

Categorical Variables: Frequency

```
In [107]: df.payment_type.value_counts()
```

```
Out[107]: Credit card    663  
         Cash           335  
         No charge       2  
         Name: payment_type, dtype: int64
```

Categorical Variables: Frequency

```
In [107]: df.payment_type.value_counts()
```

```
Out[107]: Credit card    663  
Cash                335  
No charge           2  
Name: payment_type, dtype: int64
```

```
In [108]: df.payment_type.value_counts(normalize=True)
```

```
Out[108]: Credit card    0.663  
Cash                0.335  
No charge           0.002  
Name: payment_type, dtype: float64
```


Categorical Variables: Frequency


```
In [107]: df.payment_type.value_counts()
```

```
Out[107]: Credit card    663  
Cash                335  
No charge           2  
Name: payment_type, dtype: int64
```

```
In [108]: df.payment_type.value_counts(normalize=True)
```

```
Out[108]: Credit card    0.663  
Cash                0.335  
No charge           0.002  
Name: payment_type, dtype: float64
```


```
In [109]: sns.countplot(x=df.payment_type);
```



Plotting Numeric and Categorical

Plotting Numeric and Categorical

```
In [110]: fig, ax = plt.subplots(1, 1, figsize=(12, 8))  
  
sns.barplot(x='payment_type', y='fare_amount', data=df);
```




Plotting with Hue

Plotting with Hue

```
In [111]: # Adding Another Categorical
df['day_of_week'] = df.pickup_datetime.dt.dayofweek

fig, ax = plt.subplots(1, 1, figsize=(12, 6))


sns.barplot(x='day_of_week',
            y='fare_amount',
            hue='payment_type',
            data=df);
```



Same Axis, Multiple Plots with Seaborn

Same Axis, Multiple Plots with Seaborn

```
In [112]: fig,ax = plt.subplots(1,1,figsize=(12,6))
sns.histplot(df[df.pickup_datetime.dt.hour < 12].tip_amount, label='before noon',color='blue',ax=ax);
sns.histplot(df[df.pickup_datetime.dt.hour >= 12].tip_amount, label='after noon',color='orange',ax=ax);
plt.legend(loc='best');
```



Data Exploration and Viz Review

- central tendencies: mean, median
- spread: variance, std deviation, IQR
- correlation: pearson correlation coefficient
- plotting real valued variables: histogram, scatter, regplot
- plotting categorical variables: count, bar
- plotting interactions: jointplot, pairplot

Questions?